Graded Quiz

测验, 13 个问题

13/13 分 (100%)

恭喜！您通过了！

下一项

1 / 1
分数

1.
Which approach ensures continual exploration? (Select all that apply)

- [x] Exploring starts

正确
Correct! Exploring starts guarantee that all state-action pairs are visited an infinite number of times in the limit of an infinite number of episodes.

- [ ] On-policy learning with a deterministic policy

未选择的是正确的

- [x] On-policy learning with an ϵ-soft policy

正确
Correct! ϵ-soft policies assign non-zero probabilities to all state-action pairs.

- [x] Off-Policy learning with an ϵ-soft behavior policy and a deterministic target policy

正确
Correct! ϵ-soft policies have non-zero probabilities for all actions in all states. The behavior policy is used to generate samples and should be exploratory.

- [ ] Off-Policy learning with an ϵ-soft target policy and a deterministic behavior policy

未选择的是正确的

Graded Quiz

测验, 13 个问题

1 / 1 分数

13/13 分 (100%)

2.

When can Monte Carlo methods, as defined in the course, be applied? (Select all that apply)

- [ ] When the problem is continuing and there are sequences of states, actions, and rewards

未选择的是正确的

- [ ] When the problem is continuing and there is a model that produces samples of the next state and reward

未选择的是正确的

- [ ] When the problem is episodic and there are sequences of states, actions, and rewards

正确

Correct! Well-defined returns are available in episodic tasks.

- [ ] When the problem is episodic and there is a model that produces samples of the next state and reward

正确

Correct! Well-defined returns are available in episodic tasks.

---

1 / 1 分数

3.

Which of the following learning settings are examples of off-policy learning? (Select all that apply)

- [ ] Learning about multiple policies simultaneously while following a single behavior policy

正确

Correct! Off-policy learning enables learning about multiple target policies simultaneously using a single behavior policy.

- [ ] Learning the optimal policy while continuing to explore

正确

Correct! An off-policy method with an exploratory behavior policy can assure continual exploration.

Graded Quiz

测验, 13 个问题

13/13 分 (100%)

- [x] Learning from data generated by a human expert

正确

Correct! Applications of off-policy learning include learning from data generated by a non-learning agent or human expert. The policy that is being learned (the target policy) can be different from the human expert's policy (the behavior policy).

---

1 / 1
分数

4.

If a trajectory starts at time t and ends at time T, what is its relative probability under the target policy π and the behavior policy b?

- [x] $\prod_{k=t}^{T-1} \frac{\pi(A_k \mid S_k)}{b(A_k \mid S_k)}$

正确

Correct! This is the importance sampling ratio and is used to weight returns in off-policy Monte-Carlo Policy Evaluation.

- [ ] $\sum_{k=t}^{T-1} \frac{\pi(A_k \mid S_k)}{b(A_k \mid S_k)}$
- [ ] $\frac{\pi(A_{T-1} \mid S_{T-1})}{b(A_{T-1} \mid S_{T-1})}$
- [ ] $\frac{\pi(A_t \mid S_t)}{b(A_t \mid S_t)}$

---

1 / 1
分数

5.

When is it possible to determine a policy that is greedy with respect to the value functions v_π, q_π for the policy π? (Select all that apply)

- [x] When state values v_π and a model are available

正确

Correct! With state values and a model, one can look ahead one step and see which action leads to the best combination of reward and next state.

When state values v_π are available but no model is available.

Graded Quiz

测验, 13 个问题

13/13 分 (100%)

未选择的是正确的

When action values q_π and a model are available

正确

Correct! Action values are sufficient for choosing the best action in each state.

When action values q_π are available but no model is available.

正确

Correct! Action values are sufficient for choosing the best action in each state.

---

1 / 1
分数

6.

Monte Carlo methods in Reinforcement Learning work by...

Averaging sample rewards

Planning with a model of the environment

Performing sweeps through the state set

Averaging sample returns

正确

Correct! Monte Carlo methods in Reinforcement Learning sample and average returns much like bandit methods sample and average rewards.

---

1 / 1
分数

7.

Which of the following is a requirement for using Monte Carlo policy evaluation with a behavior policy b for a target policy π?

For each state s and action a, if $b(a \mid s) > 0$ then $\pi(a \mid s) > 0$

For each state s and action a, if $\pi(a \mid s) > 0$ then $b(a \mid s) > 0$

Graded Quiz

测验, 13 个问题

13/13 分 (100%)

正确
Correct! Every action taken under π must have a non-zero probability under b.

- All actions have non-zero probabilities under π

1 / 1
分数

8.

Suppose the state s has been visited three times, with corresponding returns 8, 4, and 3. What is the current Monte Carlo estimate for the value of s?

- 3
- 15
- 5

正确
Correct! The Monte Carlo estimate for the state value is the average of sample returns observed from that state.

- 3.5

1 / 1
分数

9.

When does Monte Carlo prediction perform its first update?

- After the first time step
- When every state is visited at least once
- At the end of the first episode

正确
Correct! Monte Carlo Prediction updates value estimates at the end of an episode.

1 / 1
分数

Graded Quiz

测验, 13 个问题

13/13 分 (100%)

10.

In Monte Carlo prediction of state-values, **memory** requirements depend on (select all that apply)

- [x] The number of states

正确

Correct! Monte Carlo Prediction needs to store the estimated value for each state.

- [x] The number of possible actions in each state

未选择的是正确的

- [x] The length of episodes

正确

Correct! Monte Carlo Prediction needs to store the sequence of states and rewards. during an episode

---

1 / 1
分数

11.

For Monte Carlo Prediction of state-values, the number of **updates** at the end of an episode depends on

- [ ] The number of states
- [ ] The number of possible actions in each state
- [x] The length of the episode

正确

Correct! Monte Carlo Prediction updates the estimated value of each state visited during the episode.

---

1 / 1
分数

12.

Which approach can find an optimal deterministic policy? (select all that apply)

- [x] Exploring Starts

正确

Correct! Exploring starts ensure that every state-action pair is visited even if the policy is deterministic.

Graded Quiz
测验, 13 个问题

13/13 分 (100%)

- [x] ϵ-greedy exploration

未选择的是正确的

- [x] Off-policy learning with an ϵ-soft behavior policy and a deterministic target policy

正确

Correct! In this case, the behavior policy can maintain exploration while the target policy is deterministic.

---

1 / 1
分数

13.

In an ϵ-greedy policy over $\mathcal{A}$ actions, what is the probability of the highest valued action if there are no other actions with the same value?

- $1 - \epsilon$
- ϵ
- $1 - \epsilon + \frac{\epsilon}{\mathcal{A}}$

正确

Correct! The highest valued action still has a chance of being selected as an exploratory action.

- $\frac{\epsilon}{\mathcal{A}}$
